



## Model Performance And Fit - 1


*One should look for what is and not what he thinks should be. (Albert Einstein)*

# Model Performance and Fit: Topic introduction

In this part of the course, we will cover the following concepts:

- Implement a custom neural network to demonstrate model fit with different learning rates, epochs and batch sizes
- Understand loss functions and math behind gradient descent
- Assess and discuss methods to improve the fit of a neural network

# Warm up: Quick, Draw!



Can a neural network learn to recognize doodling?

Help teach it by adding your drawings to the [world's largest doodling data set](#), shared publicly to help with machine learning research.

Let's Draw!

- Quick, Draw! is an online game that uses neural network intelligence to guess what the drawings represent
- Visit the website to play **Quick, Draw!**
- After completing, share how many of your doodles were correctly recognized

# Module completion checklist

Objective	Complete
Summarize the role that batch size and epochs play in neural network training	
Implement a custom neural network to demonstrate model fit with different learning rates	

# Epoch

- When the entire dataset is passed through a neural network, we consider one **epoch** of the training complete
  - Since the dataset can be too big to feed into the neural network, we divide it into **several smaller batches**
  - The number of batches needed to complete the epoch equals to the number of **iterations** per epoch the algorithm needs to make

# Batch

- **Batch size** is a parameter we define while fitting the model
- It refers to the **number of observations our model takes at a time to make predictions** and update the weights in the following manner:
  - the algorithm **iterates** over one or more observations making predictions at the end of each batch
  - predictions are then compared and an error is calculated
  - this error term is then used to adjust the weights

2	NO	NO	2	11	11	81	17	331	654	Batch=1
3	NO	NO	2	9	6.5	71.5	17	131	670	
4	NO	YES	1	1	4	44	18	120	1229	
5	NO	YES	1	2	2.5	64	9	108	1454	
6	NO	YES	1	2.5	1	42.5	13	82	1518	
7	NO	YES	1	2	2	52	6	88	1518	Batch=2
8	NO	YES	2	1	3	47.5	11	148	1362	
9	NO	NO	2	1	5	51	17	68	891	
10	NO	NO	1	2	8.5	46	25	54	768	
11	NO	YES	1	2	6	50	15	41	1280	
.....										
727	NO	YES	3	2.5	3	87	22	9	432	Batch=n
728	NO	YES	2	3	1	60	21	247	1867	
729	NO	YES	2	3	1	60	9	644	2451	
730	NO	NO	2	3	2	75	6	159	1182	
731	NO	NO	1	4	1	51	24	364	1432	

# Batch order

- Batches are selected randomly during each epoch, unless we specify otherwise
  - Most problems are agnostic to the order of observations in batches and the order of batches
  - Some models, like time series problems, do need to have observations appear in order and need batch options set accordingly

5	NO	YES	1	1.5	2	70	12	15	416	Batch=1
10	NO	YES	1	2	6	50	15	41	1280	
22	NO	YES	1	3	4	40	27	64	1466	
26	NO	YES	3	11.5	11.5	94	7	69	1538	
31	NO	YES	1	2	5	42	22	75	1468	
33	NO	YES	2	10	8	93	19.5	81	1365	Batch=2
35	NO	YES	2	4	2	56	15	83	1844	
36	NO	YES	1	2.5	7	49	13	86	1330	
37	NO	YES	2	6.3	6.7	85.6	21.12	87	1009	
40	NO	YES	1	2	8	42.5	24	89	2147	
.....										
703	NO	NO	1	23	23	58	19	2512	5883	Batch=n
713	NO	NO	1	19.5	19.5	49	9	2622	4807	
719	NO	NO	1	18	18	52	11	2795	4665	
720	NO	NO	1	20	20	49	13	2795	5325	
730	NO	NO	1	13.5	13.5	23.5	17	3252	3605	

# Batch size implications

- Batch size plays a big role in the time it takes to train the model and how well the model is doing
- With a **small** batch size, expect:
  - a longer train time
  - higher accuracy (which may lead to model overfitting)
  - it easily fits in computer memory
- With a **medium** batch size, expect:
  - a balance between training time and accuracy
  - better generalization to new data
  - it fits in memory
- With a **large** batch size, expect:
  - a faster training time, especially if training is done in parallel
  - it may generalize to new data poorly
  - it could be too big to fit into memory



# Batch size implications (cont'd)

- A single batch can:
  - consist of a single observation (`batch_size = 1`), known as **stochastic gradient descent**
  - be as large as your dataset (`batch_size = num_observations`), known as **batch gradient descent**
  - be any value between 1 and the total `num_observations`, known as **mini-batch gradient descent**
- The default value for a batch size in TensorFlow and many other neural network frameworks is 32

# Batch size calculations

- `num_observations / batch_size` will provide the total number of resulting batches
- If the `batch_size` doesn't evenly divide into `num_observations`, the last batch will be smaller than the rest
  - If the model is sensitive to batches having the same outer dimension, set the `drop_remainder` argument to `True` to drop the smaller batch

# Example: batch size, iterations, and epochs math

- Let's say you have a dataset with 1000 observations and you choose to set the batch size to 50 and the epochs to 200:
  - `num_observations = 1000`
  - `batch_size = 50`
  - `epochs = 200`
- This means that the entire dataset will be split into 20 batches, each with 50 observations:
  - `num_batches = num_observations / batch_size = 20`
- This also means that 1 epoch will involve 20 batches or 20 updates in the model
  - `iterations = num_batches = 20`
- With 200 epochs, the model passes through the entire dataset 200 times, which means
  - `total_iterations = epochs*iterations = 4000`

# Module completion checklist

Objective	Complete
Summarize the role that batch size and epochs play in neural network training	✓
Implement a custom neural network	

# Loading packages

- Let's load the packages we will be using:

```
# Helper packages.
import os
import pickle
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import math
import seaborn as sns

# Scikit-learn packages.
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
from sklearn import metrics
```

```
# TensorFlow and supporting packages.
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.models import load_model
from tensorflow.keras.optimizers import Adam
from tensorflow.keras import Sequential
from tensorflow.keras.layers import Dense
```

# Directory settings

- Let's start by encoding the directory structure into `variables` in order to maximize the efficiency of your workflow
- Let the `data_dir` be the variable corresponding to your `data` folder

```
from pathlib import Path
home_dir = Path(".").resolve()
main_dir = home_dir.parent.parent
print(main_dir)
```

```
data_dir = str(main_dir) + "/data"
print(data_dir)
```

# Load the data

- The `credit_card_data` dataset contains information about credit card **defaulters**
- Our goal is to predict if the customer will **default on a credit card payment**

```
credit_card = pd.read_csv(str(data_dir) + "/credit_card_data.csv")  
print(credit_card.head())
```

	ID	LIMIT_BAL	SEX	...	PAY_AMT5	PAY_AMT6	default_payment_next_month
0	1	20000	2	...	0	0	1
1	2	120000	2	...	0	2000	1
2	3	90000	2	...	1000	5000	0
3	4	50000	2	...	1069	1000	0
4	5	50000	1	...	689	679	0

[5 rows x 25 columns]

# Data preparation

- We need to make sure our data is in a suitable form to run through a neural network, which is why we must:
  - check the data for **NAs**
  - **transform the data to numeric values** (*if it's categorical, make sure the data is encoded*)
  - split data into train, test, and validation
  - normalize data
  - examine the target variable imbalance

Remember that the order of operations matters!

- Ideally, all data transformations should happen after the data has been split
- In this instance, we will check for NAs and encode categorical variables before the split since this will not significantly affect the results and will keep our code more concise



# Data prep: convenience function

- We wrote a time-saving function to perform all of the cleaning and split steps on the credit card dataset at once!

```
def data_prep(df):  
  
    # Fill missing values with mean  
    df = df.fillna(df.mean()['BILL_AMT1'])  
    # Drop an unnecessary identifier column.  
    df = df.drop('ID', axis = 1)  
  
    # Convert 'sex' into dummy variables.  
    sex = pd.get_dummies(df['SEX'], prefix = 'sex', drop_first = True)  
    # Convert 'education' into dummy variables.  
    education = pd.get_dummies(df['EDUCATION'], prefix = 'education', drop_first = True)  
    # Convert 'marriage' into dummy variables.  
    marriage = pd.get_dummies(df['MARRIAGE'], prefix = 'marriage', drop_first = True)  
    # Drop `sex`, `education`, `marriage` from the data.  
    df.drop(['SEX', 'EDUCATION', 'MARRIAGE'], axis = 1, inplace = True)  
  
    # Concatenate `sex`, `education`, `marriage` dummies to our dataset.  
    df = pd.concat([df, sex, education, marriage], axis=1)  
  
    # Separate predictors from data.  
    X = df.drop(['default_payment_next_month'], axis=1)  
  
    # Separate target from data.  
    y = df['default_payment_next_month']
```

# Data prep: convenience function (cont'd)

```
# Set the seed to 1.
np.random.seed(1)

# Split data into train, test, and validation set, use a 70 - 15 - 15 split.
# First split data into train-test with 70% for train and 30% for test.
X_train, X_test, y_train, y_test = train_test_split(X.values,
                                                    y,
                                                    test_size = .3,
                                                    random_state = 1)

# Then split the test data into two halves: test and validation.
X_test, X_val, y_test, y_val = train_test_split(X_test,
                                                y_test,
                                                test_size = .5,
                                                random_state = 1)

print("Train shape:", X_train.shape, "Test shape:", X_test.shape, "Val shape:", X_val.shape)

# Transforms features by scaling each feature to a given range.
# The default is the range between 0 and 1.
min_max_scaler = preprocessing.MinMaxScaler()
X_train_scaled = min_max_scaler.fit_transform(X_train)
X_test_scaled = min_max_scaler.transform(X_test)
X_val_scaled = min_max_scaler.transform(X_val)

return X_train_scaled, X_test_scaled, X_val_scaled, y_train, y_test, y_val
```

# Data prep

```
X_train_scaled, X_test_scaled, X_val_scaled, y_train, y_test, y_val = data_prep(credit_card)
```

```
Train shape: (21000, 30) Test shape: (4500, 30) Val shape: (4500, 30)
```

# Define and compile a sequential model

- Let's create a convenience function to define and compile the model with an input layer, two hidden layers, and an output layer

```
def create_model(lr = .01):  
    # Let's set the seed so that we can reproduce the results.  
    tf.random.set_seed(1)  
    opt = Adam(learning_rate = lr) # <- set optimizer  
  
    model = Sequential([  
        Dense(32, activation='relu', input_dim=30), #<- set input and 1st hidden layer  
        Dense(32, activation='relu'),               #<- set 2nd hidden layer  
        Dense(1, activation='sigmoid')              #<- set output layer  
    ])  
  
    model.compile(optimizer = opt,                #<- set optimizer  
                  loss='binary_crossentropy',    #<- set loss function to binary_crossentropy  
                  metrics=['accuracy'])          #<- set performance metric  
    return model
```

- We will create models with different numbers of learning rates to compare how those parameters affect loss and accuracy
- We will use the default `batch_size` - 32 in this experiment

# Default learning rate

- We have set the learning rate to default to 0.01 in the `create_model()` function, so we don't have to specify it explicitly

```
lr_default = create_model().fit(X_train_scaled, y_train,  
                                epochs = 25,  
                                validation_data=(X_val_scaled, y_val))
```

```
Epoch 1/25  
657/657 [=====] - 1s 1ms/step - loss: 0.4889 - accuracy: 0.7938 -  
val_loss: 0.4522 - val_accuracy: 0.8196  
Epoch 2/25  
657/657 [=====] - 1s 942us/step - loss: 0.4496 - accuracy: 0.8161 -  
val_loss: 0.4504 - val_accuracy: 0.8162  
...  
Epoch 24/25  
657/657 [=====] - 1s 852us/step - loss: 0.4372 - accuracy: 0.8169 -  
val_loss: 0.4337 - val_accuracy: 0.8202  
Epoch 25/25  
657/657 [=====] - 1s 827us/step - loss: 0.4294 - accuracy: 0.8236 -  
val_loss: 0.4355 - val_accuracy: 0.8216
```

# High learning rate

- Let's set the learning rate to a very high number like 0.75

```
# Set learning rate to 0.75.
lr_high = create_model(lr = .75).fit(X_train_scaled, y_train,
                                     epochs = 25,
                                     validation_data=(X_val_scaled, y_val))
```

```
Epoch 1/25
657/657 [=====] - 1s 930us/step - loss: 1.5291 - accuracy: 0.7780 -
val_loss: 0.5284 - val_accuracy: 0.7798
Epoch 2/25
657/657 [=====] - 1s 831us/step - loss: 0.5381 - accuracy: 0.7773 -
val_loss: 0.5272 - val_accuracy: 0.7798
...
Epoch 24/25
657/657 [=====] - 1s 840us/step - loss: 0.5459 - accuracy: 0.7751 -
val_loss: 0.5336 - val_accuracy: 0.7798
Epoch 25/25
657/657 [=====] - 1s 859us/step - loss: 0.5362 - accuracy: 0.7824 -
val_loss: 0.5274 - val_accuracy: 0.7798
```

# Low learning rate

- Let's set the learning rate to a very low number like 0.0001
  - Let's also increase the number of epochs here to 50 since the learning rate is low

```
lr_low = create_model(lr=.0001).fit(X_train_scaled, y_train,  
                                     epochs = 50, #<- increase the number of epochs  
                                     validation_data=(X_val_scaled, y_val))
```

```
Epoch 1/50  
657/657 [=====] - 1s 952us/step - loss: 0.5974 - accuracy: 0.7577 -  
val_loss: 0.5131 - val_accuracy: 0.7798  
Epoch 2/50  
657/657 [=====] - 1s 841us/step - loss: 0.5114 - accuracy: 0.7773 -  
val_loss: 0.4983 - val_accuracy: 0.7800  
...  
Epoch 49/50  
657/657 [=====] - 1s 832us/step - loss: 0.4391 - accuracy: 0.8223 -  
val_loss: 0.4392 - val_accuracy: 0.8227  
Epoch 50/50  
657/657 [=====] - 1s 830us/step - loss: 0.4328 - accuracy: 0.8254 -  
val_loss: 0.4393 - val_accuracy: 0.8233
```

# Visualize results for learning rates

- Let's create a dataframe with the loss and accuracy for training and validation data along with their corresponding epochs and learning rates

```
learn_rates = []

for exp, result in zip([lr_default, lr_low, lr_high], [".01", ".0001", ".75"]):

    df = pd.DataFrame.from_dict(exp.history)
    df['epoch'] = df.index.values
    df['Learning Rate'] = result


    learn_rates.append(df)

df_learning = pd.concat(learn_rates)
df_learning['Learning Rate'] = df_learning['Learning Rate'].astype('str')
```




# Visualize results for learning rates (cont'd)

```
sns.lineplot(x='epoch', y='val_loss', hue='Learning Rate', data=df_learning)
```



# Visualize results for learning rates (cont'd)

```
sns.lineplot(x='epoch', y='val_accuracy', hue='Learning Rate', data=df_learning)
```



- We obtain the best results when the learning rate is set to 0.01 and the model seems to clearly underfit when the learning rate is high

# Knowledge check



# Module completion checklist

Objective	Complete
Summarize the role that batch size and epochs play in neural network training	✓
Implement a custom neural network to demonstrate model fit with different learning rates	✓

# Congratulations on completing this module!

You are now ready to try tasks 1-7 in the Exercise for this topic

